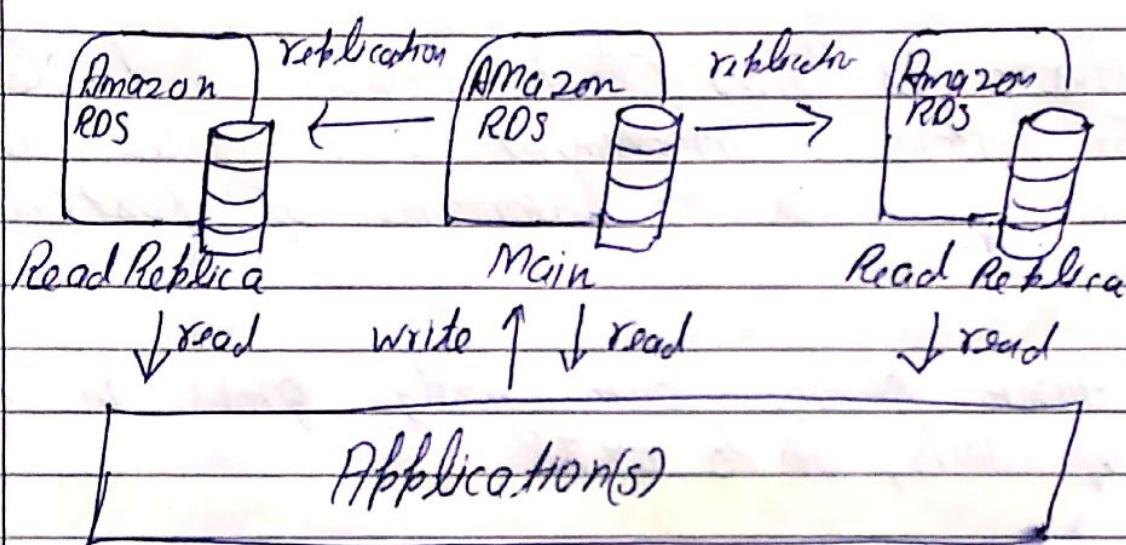


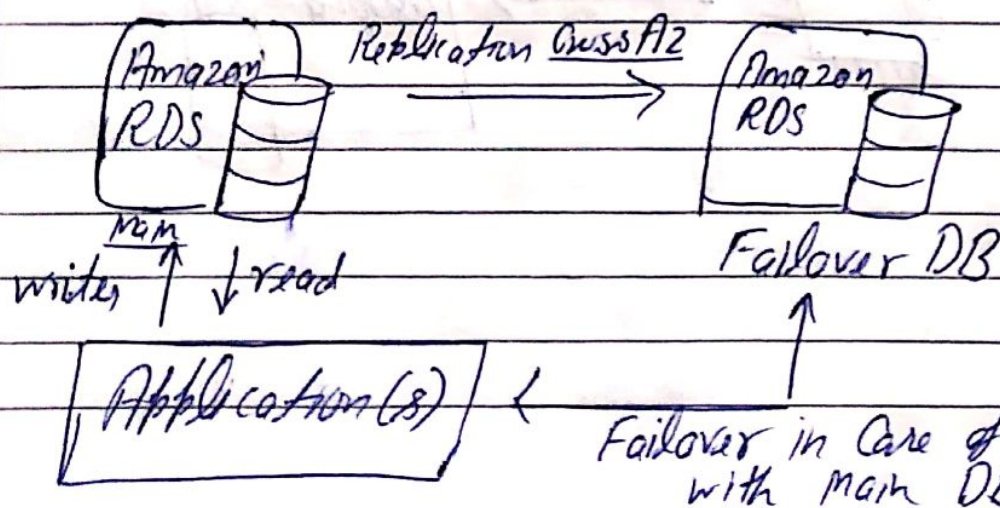
Day 14/100 →

RDS Deployments: Read Replicas, Multi-AZ• Read Replicas

- Scale the read workload of your DB
- Can Create up to 5 Read Replicas
- Data is only written to the main DB.

• Multi-AZ:

- Failover in case of AZ outage (High availability)
- Data is only read/write to the main database
- Can only have 1 other AZ as failover

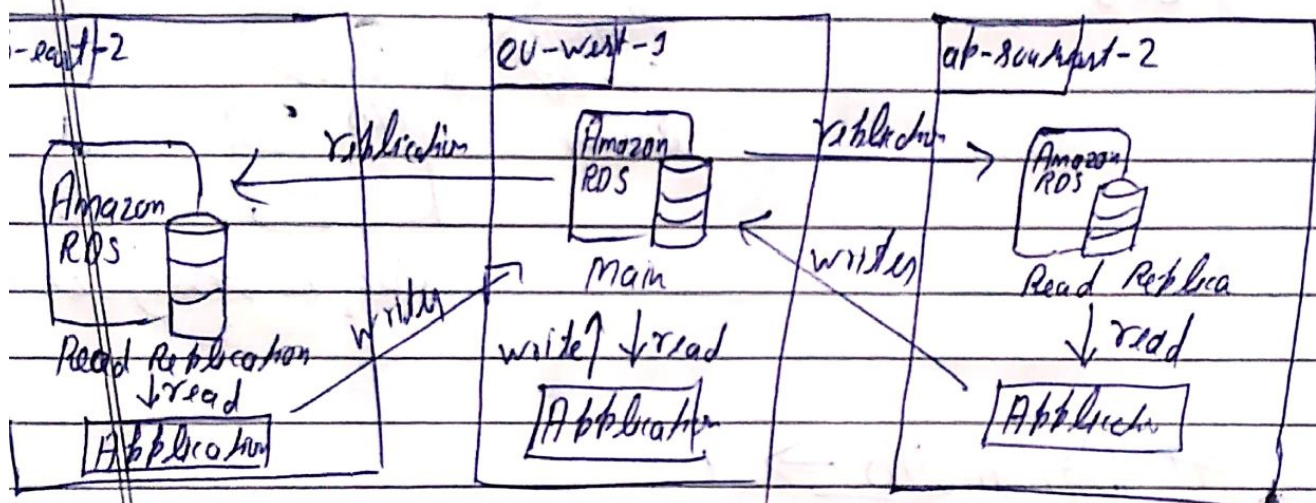


Failover in case of issues with main DB

RDS Deployments : Multi-Regions

• Multi-Region (Read Replicas)

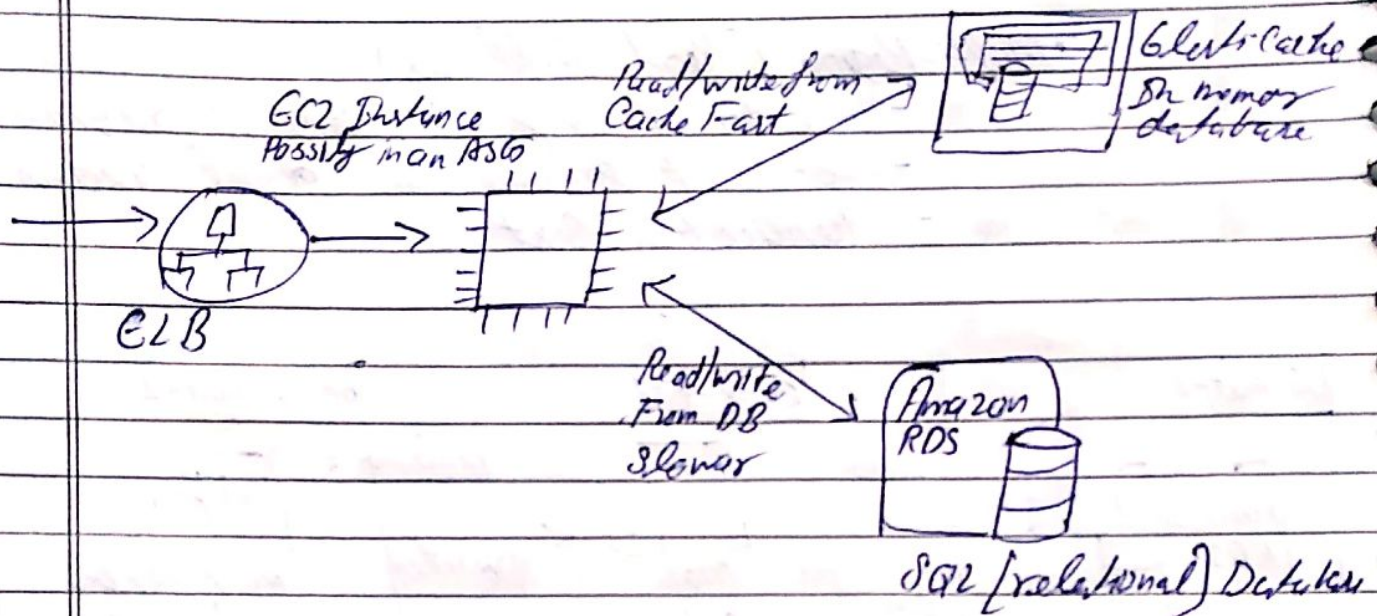
- Disaster recovery in case of region issue.
- Local performance for global reads
- Replication Cost



Amazon ElastiCache Overview

- The same way RDS is to get managed Relational Databases.
- ElastiCache is to get managed Redis or memcached
- Caches are in-memory databases with high performance, low latency
- Helps reduce of load off database for read intensive workloads.
- AWS takes care of OS maintenance/patching, optimization, setup Configuration, monitoring, failure recovery and backups

ElasticCache Solution Architecture - Cache

# DynamoDB →

- Fully Managed Highly available with replication across 3 AZ
- NoSQL database - not a relational DB
- Scales to massive workloads, distributed "serverless" database
- millions of requests per second, trillions of rows, 100s TB of storage

- Fast and Consistent in performance.
- Single-digit millisecond Latency - low latency retrieval
- Integrated with IAM for security, authentication and administration
- Low Cost and auto scaling capabilities

DynamoDB - type of data

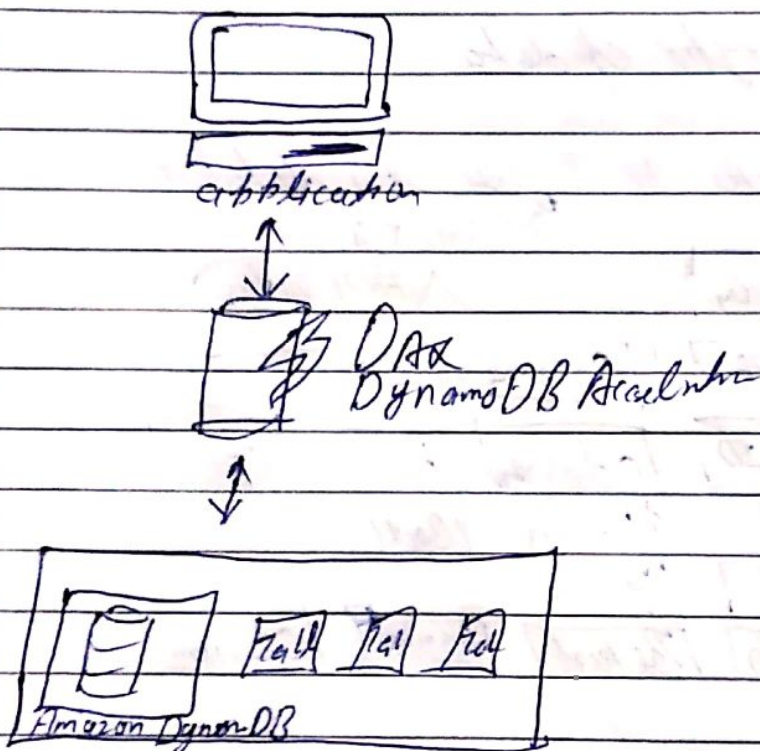
Primary key - DynamoDB is a key value database
Products

Partition Key, Sort Key		Distributed		
Product ID	Type	Schema is defined Partition		
1	Book ID	Odyssey	Home	1871
2	Album ID	Garage	Back	
2	Album ID, Track ID	Park Ave	NO. 9	
3	Movie ID	The Kid	Drama Comedy	Chetlin

DynamoDB Acceleration - DAX

- Fully Managed in-memory Cache for DynamoDB
- 10x performance improvement - Single - digit millisecond latency to microseconds latency - when accessing your DynamoDB tables

- Secure, highly scalable & highly available
- Difference with ElastiCache at the CCP level: DAX is only used for and is integrated with DynamoDB, while ElastiCache can be used for other databases.



Redshift Overview

- OLTP
Online
Transaction
Processing
- Redshift is based on PostgreSQL, but it's not used for OLTP.
 - It's OLAP - Online analytical processing (analytics & data warehousing)

- Load data Once every hour, not every second
- 10x better performance than other data warehouses, scale to PBs of data.
- Columnar Storage of Data Instead of row based
- Massive Parallel Query Execution (MPQ), highly available
- Pay as you go based on the instance provisioned
- Has SQL interface for performing the queries
- BI tools such as AWS QuickSight or Tableau integrate with it

Amazon EMR

- EMR stands for "Elastic Map Reduce"
- EMR stands helps creating Hadoop clusters (Big Data) to analyze and process vast amount of data
- The clusters can be made of hundreds of EC2 instances
- Also supports apache Spark, HBase, Presto, Flink ---

- EMR takes care of all provisioning & Configuration
- Auto-Scaling and integrated with Spot Instances.
- Use Cases: data processing, machine learning, web industries, big data

Athena

- Fully Serverless database with SQL capabilities
 - Used to query data in S3
 - Pay per query
 - Output results back to S3
 - Secured through IAM
- Use Case: One-time SQL queries, serverless queries on S3, log analytics.

AMAZON Quick Sight

- Serverless machine learning-powered business intelligence service to create interactive dashboards
- Fast, automatically scalable, embeddable, with per-session pricing

• Use Cases:

- Business Analytics
- Building Visualizations
- Perform ad-hoc analysis
- Get business insights using data

• Integrated with RDS, Aurora, Athena, Redshift, S3---